

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/2489562>

Chord: A Scalable Peer-to-Peer Lookup Service for Internet Applications

Article in *ACM SIGCOMM Computer Communication Review* · December 2001

DOI: 10.1145/964723.383071 · Source: CiteSeer

CITATIONS

8,402

READS

1,891

5 authors, including:



Ion Stoica

University of California, Berkeley

432 PUBLICATIONS 79,532 CITATIONS

[SEE PROFILE](#)



David Ron Karger

Massachusetts Institute of Technology

445 PUBLICATIONS 56,077 CITATIONS

[SEE PROFILE](#)

Some of the authors of this publication are also working on these related projects:



Spinal codes [View project](#)



X-Trace [View project](#)

Chord: A Scalable Peer-to-Peer Lookup Service for Internet Applications

Ion Stoica

David Liben-Nowell

M. Frans Kaashoek

Robert Morris


David R. Karger

Frank Dabek

Hari Balakrishnan

CS856


Nabeel Ahmed






Outline

- P2Ps as Lookup Services
- Related Work
- Chord System Model
- Chord Protocol Description
- Simulation Results
- Current Status and Issues
- Extensions of Chord
- References
- Discussion




A P2P Lookup Service?

- P2P system:
 - Data items spread over a large number of nodes
 - Which node stores which data item?
 - A lookup mechanism needed
 - Centralized directory -> bottleneck/single point of failure
 - Query Flooding -> scalability concerns
 - Need more structure!
- Solution: *Chord* (a distributed lookup protocol)
- Chord supports only one operation: *given key, maps key on to a node*




Related Work

- Unstructured Peer-to-Peer Systems
 - Freenet
 - KaZaa/Napster
 - Gnutella
- Structured Peer-to-Peer Systems
 - CAN
 - OceanStore (Tapestry)
 - Pastry
 - Kademlia, Viceroy etc..
- To many routing structures? How to compare?




Related Work (Contd..)

- Routing Geometry: *“Manner in which neighbors and routes are chosen” Gummadi et al.[6]*
- Classify Routing Geometries:
 - Tree → *PRR, Tapestry, Globe system, TOPLUS*
 - Hypercube → *CAN,*
 - Butterfly → *Viceroy*
 - Ring → *Chord*
 - XOR → *Kademlia*
 - Hybrid → *Pastry (Tree/Ring)*
 - *Maybe more....*
- *Compare degree of flexibility in routing geometries*
 - Neighbor Selection
 - Route Selection
- *Comparative discussion later.....*




Chord System Model

- Design Objectives:
 - Load Balance: Distributed hash function spreads keys evenly over the nodes
 - Decentralization: Fully distributed
 - Scalability: Lookup grows as a log of number of nodes
 - Availability: Automatically adjusts internal tables to reflect changes.
 - Flexible Naming: No constraints on key structure.
- Example Applications:
 - Co-operative Mirroring
 - Time-shared storage
 - Distributed indexes
 - Large-Scale combinatorial search



Chord Protocol

- Assumption: Communication in underlying network is both symmetric and transitive.
- Assigns keys to nodes using *consistent hashing*
- Uses *logical ring* geometry to manage identifier space (identifier circle)
- Utilizes (*sequential*) *successor/predecessor pointers* to connect nodes on ring
- Distributes routing table among nodes (*Finger pointers*)



Consistent Hashing


- Properties:
 - Minimal Disruption: require minimal key movement on node joins/leaves
 - Load Balancing: distribute keys equally across over nodes

Theorem: For any set of N nodes and K keys, with *high probability*:

- 1) Each node is responsible for at most $(1+e)K/N$ keys.
- 2) When an $(N+1)$ st node joins or leaves the network, responsibility for $O(K/N)$ keys changes hands.

$$e = O(\log N)$$


Consistent Hashing (Contd..)




- Consistent hashing function assigns each node and key an m -bit identifier using SHA-1 base hash function (*160-bits truncated to m*).
- Node's IP address is hashed.
- Identifiers are ordered on a identifier circle modulo 2^m called a chord ring.
- $successor(k)$ = first node whose identifier is \geq identifier of k in identifier space



Example Chord Ring





$m = 6$
10 nodes




Lookups in Chord

- Two techniques:
 - Simple-Key Location scheme:
 - State-maintenance $O(1)$ [no finger table]
 - Lookup-time $O(N)$ [follow successor pointers]
 - Scalable-Key Location scheme:
 - State-maintenance $O(\log N)$ [finger table]
 - Lookup-time $O(\log N)$ [follow finger pointers]

Simple Key Location Scheme





Scalable Key Lookup Scheme



- Finger Pointers
 - $n.\text{finger}[i] = \text{successor}(n + 2^{i-1})$
 - Each node knows more about portion of circle close to it!
- Query the finger-node that is nearest predecessor of key (closest preceding finger)
- Recursive querying till immediate predecessor p of key found
- Return $p.\text{successor}$

Scalable Lookup Scheme: Finger Table





Finger Table for N8

N8+1	N14
N8+2	N14
N8+4	N14
N8+8	N21
N8+16	N32
N8+32	N42

finger [k] = first node that succeeds $(n+2^{k-1}) \bmod 2^m$

Scalable Lookup Scheme





What about Churn?

- **Churn:** Term used for dynamic membership changes
- Problems related to Churn:
 - Re-delegation of key-storage responsibility
 - Updation of finger tables for routing
- Need to support:
 - Concurrent Node Joins/Leaves (Stabilization)
 - Fault-tolerance and Replication (Robustness)



Node Joins

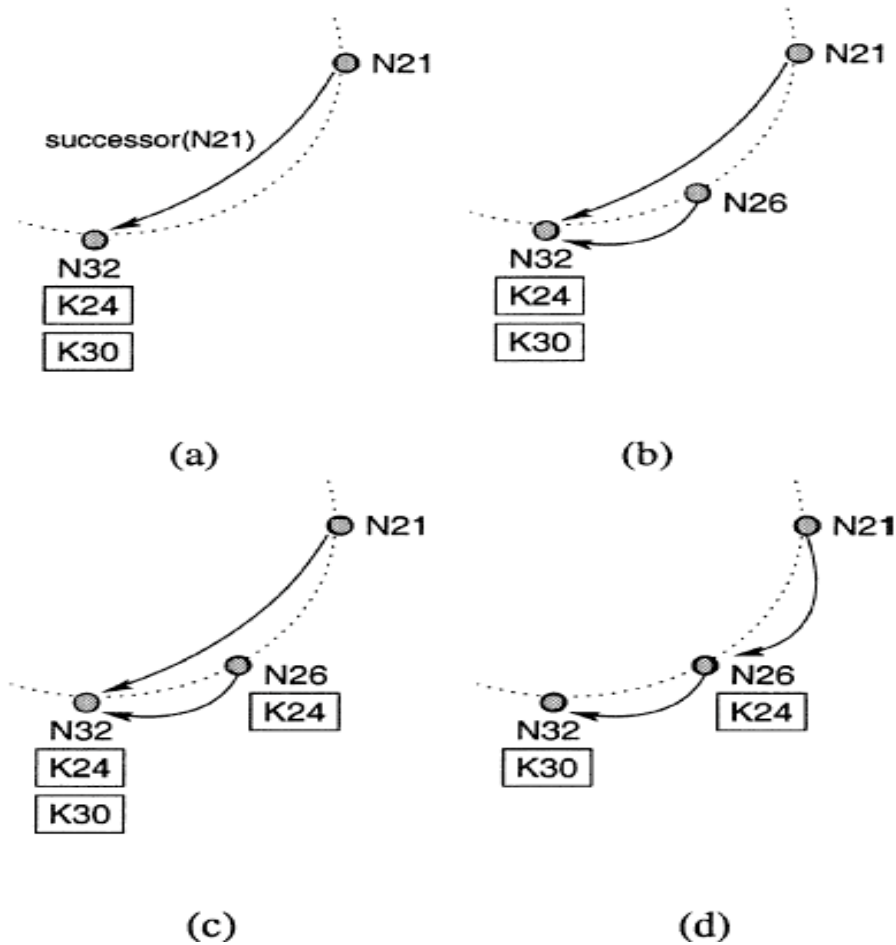
- New node B learns of at least one existing node A via external means
- B asks A to lookup its finger-table information
 - Given B's hash-id b , A does lookup for $B.finger[i] = \text{successor}(b + 2^{i-1})$ if interval not already included in $finger[i-1]$
 - B stores all finger information and sets up pred/succ pointers
- Updation of finger table required at certain existing nodes
- Key movement is done from $\text{successor}(b)$ to b



Concurrent Joins/Leaves

- Problem: Join operation difficult to run for concurrent joins/leaves in large networks
- Solution: Use a **stabilization** protocol that runs periodically to guard against inconsistency
- Each node periodically runs stabilization protocol
 - Check consistency of succ. pointer <basic stabilization>
 - Check consistency of finger pointers <fix_fingers>
 - Check consistency of pred. pointer <check_predecessor>
- Note:
 - Stabilization protocol guarantees to add nodes in a fashion to preserve reachability
 - Incorrect finger pointers may only increase latency, but incorrect successor pointers may cause lookup failure!


Modified Node Join



Fault-tolerance and Replication




- Fault-tolerance:
 - Maintain successor invariant
 - Each node keeps track of r successors
 - If $r = O(\log(N))$, then lookups succeed with high probability despite a failure probability of $\frac{1}{2}$
- Replication:
 - Supports replication by storing each item at some k of these r successor nodes



Voluntary Node Departures



- Can be treated as node failures
- Two possible enhancements
 - Leaving node may transfers all its keys to its successor
 - Leaving node may notify its predecessor and successor about each other so that they can update their links



Simulation Results




- Iterative implementation
- 10,000 nodes
- No. of keys range from 10^5 to 10^6
- Presented results:
 - Load Balance
 - Path Length
 - Lookups during stabilization
- Comparative discussion on DHTs

Load Balance




**Drastic Variation
in Key
Allocation:
Poor Load
Balance**

Path Length



Lookups during Stabilization



Comparative Discussion on DHTs



- Comparison metrics: (*degree of flexibility*) Gummadi et. al [6]
 - Static Resilience: Ability to route successfully w/out recovery
 - Path Latency: Average end-to-end latency for a lookup
 - Local Convergence: Property that 2 messages for same location converge at a node near the two sources
- From study, [6] conclude ring-structure performs the best!

property	tree	hypercube	ring	butterfly	xor	hybrid
Neighbor Selection	$n^{\log n/2}$	1	$n^{\log n/2}$	1	$n^{\log n/2}$	$n^{\log n/2}$
Route Selection (optimal paths)	1	$c_1(\log n)$	$c_1(\log n)$	1	1	1
Route Selection (non-optimal paths)	-	-	$2c_2(\log n)$	-	$c_2(\log n)$	$c_2(\log n)$
Natural support for sequential neighbors?	no	no	yes	no	no	Default routing: no Fallback routing: yes




Current Status

- Is actively being investigated as project IRIS:
 - Infrastructure for Resilient Internet Systems (<http://project-iris.com/>)
 - Government funded project active since 2002 (\$12M)
 - Goal: “*develop novel decentralized infrastructure based on distributed hash-tables that enable a new generation of large-scale distributed applications*”.
- Has been used in:
 - General-purpose DHASH layer for various applications
 - DDNS (Distributed DNS)
 - CFS (Wide-area Co-operative File System for distributed read-only storage)
 - Ivy (peer-to-peer read/write file-system)
 - Internet Indirection Infrastructure (I3)




Still many issues...

- Security considerations: (many possible attacks beyond data integrity)
 - Routing attacks: incorrect lookups/updates/partitions
 - Storage & Retrieval attacks: *denial-of-service/data*
 - Other misc. attacks: inconsistent behavior, overload, etc.
- Performance considerations:
 - No consideration of underlying routing topology (locality properties)
 - No consideration of underlying network traffic/congestion condition
 - Bound on lookups still not good enough for some applications
 - E.g. Failure of DDNS since 8-orders of magnitude worse than conv. DNS
- Application-Specific considerations:
 - Each application requires its own set of access functions in the DHT
 - Lack of sophisticated API for supporting such applications
 - E.g. DHASH API is too basic to support sophisticated functionality
 - Support only for DHT as library vs. as a service
- *And many more...*




Extensions of Chord

- Hierarchical Chord (Crescendo)
 - “*Canon*” generic transformation applied to create hierarchy structure on any flat DHT.
 - Each domain/sub-domain in hierarchy is represented by a ring
 - Larger domains consist of *merged* ring of smaller domains
 - Is this adequate for *locality properties*?



Hierarchy of Domains





Merging two Chord Rings


Extensions of Chord (Contd..)



- Internet Indirection Infrastructure (*i3*)
 - Combines Chord's lookup with forwarding
 - Receiver inserts trigger (Id, R) into ring
 - Sender sends data to receiver's Id
- Supports:
 - Mobility with location privacy (ROAM)
 - Multicast/ Anycast
 - Service-composition



References



- [1] E. Sit and R. Morris, *Security Considerations for Peer-to-Peer Distributed Hash Tables*, In the proceedings of the First International Workshop on Peer-to-Peer Systems (IPTPS '02), March, 2002; Cambridge, MA
- [2] F. Dabek, E. Brunskill, F. Kaashoek, D. Karger, R. Morris, I. Stoica, and H. Balakrishnan, *Building Peer-to-Peer Systems With Chord, a Distributed Lookup Service*, Proceedings of the 8th Workshop on Hot Topics in Operating Systems (HotOS-VIII), May 2001
- [3] R. Cox, A. Muthitacharoen, R. Morris, *Serving DNS using a Peer-to-Peer Lookup Service*, In the proceedings of the First International Workshop on Peer-to-Peer Systems (IPTPS '02), March, 2002; Cambridge, MA
- [4] B. Karp, S. Ratnasamy, S. Rhea, and S. Shenker. Spurring Adoption of DHTs with OpenHash, a Public DHT Service, *In Proceedings of the 3rd International Workshop on Peer-to-Peer Systems (IPTPS '04)*, February 2004
- [5] Ganesan, Prasanna; Gummadi, Krishna; Garcia-Molina, Hector. Canon in G Major: Designing DHTs with Hierarchical Structure, Proc. International Conference on Distributed Computing Systems (ICDCS) 2004 .
- [6] K. Gummadi, R. Gummadi, S. Gribble, S. Ratnasamy, S. Shenker, I. Stoica, The Impact of DHT Routing Geometry on Resilience Proximity, *In Proceedings of ACM SIGCOMM 2003*
- [7] I. Stoica, D. Adkins, S. Zhuang, S. Shenker, S. Surana, "Internet Indirection Infrastructure," *Proceedings of ACM SIGCOMM*, August, 2002
- [8] Host Mobility using an Internet Indirection Infrastructure, *First International Conference on Mobile Systems, Applications, and Services (ACM/USENIX Mobisys)*, May, 2003



Discussion

- Chord could still suffer from potential network partitioning problems
 - How to enforce stricter guarantees on robustness with minimal additional overhead?
- How scalable is the stabilization protocol?
 - Is there a stabilization rate that is suitable for all deployments?
 - How do we balance consistency and network overhead?
- Utilize *caching* on search path for performance?
 - Improve performance for popular DHT lookups (*hay*)
 - Cache coherency problems?
- Performance and Security seem to be at direct odds with each other
 - Can we provide a solution that supports both?
- What is a better approach, DHTs as a library? Or as a service?
- How can we incorporate query models beyond exact-matches?
- What adoption incentives do DHTs need to provide?